CS 747, Autumn 2020: Week 0, Lecture 1

Shivaram Kalyanakrishnan

Department of Computer Science and Engineering
Indian Institute of Technology Bombay

Autumn 2020

What is this Course About?

- Are you wondering: "Should I register for this course or not?"

What is this Course About?

- Are you wondering: "Should I register for this course or not?"

- How can you make the right decision, especially since you are uncertain about
  - how much time you'll be able to devote to it;
  - whether you have sufficient background;
  - the course's relevance to your future plans;
  - whether the topics will be interesting;
  - whether the videos will induce sleep;
  - all the same questions about other courses you could do instead?!

What is this Course About?

- Are you wondering: "Should I register for this course or not?"
- How can you make the right decision, especially since you are uncertain about
  - how much time you'll be able to devote to it;
  - whether you have sufficient background;
  - the course's relevance to your future plans;
  - whether the topics will be interesting;
  - whether the videos will induce sleep;
  - all the same questions about other courses you could do instead?!
- Life is full of uncertainty, and yet life is full of decisions!

What is this Course About?

- Are you wondering: "Should I register for this course or not?"
- How can you make the right decision, especially since you are uncertain about
  - how much time you'll be able to devote to it;
  - whether you have sufficient background;
  - the course's relevance to your future plans;
  - whether the topics will be interesting;
  - whether the videos will induce sleep;
  - all the same questions about other courses you could do instead?!
- Life is full of uncertainty, and yet life is full of decisions!
- This course—"Foundations of Intelligent and Learning Agents"—is about decision making in the face of uncertainty.

What is this Course About?

- Are you wondering: "Should I register for this course or not?"
- How can you make the right decision, especially since you are uncertain about
  - how much time you'll be able to devote to it;
  - whether you have sufficient background;
  - the course's relevance to your future plans;
  - whether the topics will be interesting;
  - whether the videos will induce sleep;
  - all the same questions about other courses you could do instead?!
- Life is full of uncertainty, and yet life is full of decisions!
- This course—"Foundations of Intelligent and Learning Agents"—is about decision making in the face of uncertainty.
- If you want to decide in a principled way whether to take this course,

What is this Course About?

- Are you wondering: "Should I register for this course or not?"
- How can you make the right decision, especially since you are uncertain about
  - how much time you'll be able to devote to it;
  - whether you have sufficient background;
  - the course's relevance to your future plans;
  - whether the topics will be interesting;
  - whether the videos will induce sleep;
  - all the same questions about other courses you could do instead?!
- Life is full of uncertainty, and yet life is full of decisions!
- This course—"Foundations of Intelligent and Learning Agents"—is about decision making in the face of uncertainty.
- If you want to decide in a principled way whether to take this course, you need to first take this course!

Main Topics

- **Multi-armed bandits**
- **Markov Decision Problems**
- **Reinforcement learning**
- **Multi-agent systems/learning**

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
- **Reinforcement learning**
- **Multi-agent systems/learning**

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**

- **Multi-agent systems/learning**

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**
  - Decision making in the presence of other decision-makers.

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**
  - Decision making in the presence of other decision-makers.

- Applications in game playing, robotics and control, planning and scheduling on-line advertising, autonomous navigation, etc.

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**
  - Decision making in the presence of other decision-makers.

- Applications in game playing, robotics and control, planning and scheduling on-line advertising, autonomous navigation, etc.

  We will consider some case studies as a part of the course.

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**
  - Decision making in the presence of other decision-makers.

- Applications in game playing, robotics and control, planning and scheduling on-line advertising, autonomous navigation, etc.

  We will consider some case studies as a part of the course.

- Previous years’ course pages have weekly schedule, readings, references, assignments, assessments. Take a look.

Main Topics

- **Multi-armed bandits**
  - The “explore or exploit” tradeoff.
- **Markov Decision Problems**
  - Sequential decision making.
- **Reinforcement learning**
  - Learning by trial and error, reward and punishment, to optimise long-term gain.
- **Multi-agent systems/learning**
  - Decision making in the presence of other decision-makers.

- Applications in game playing, robotics and control, planning and scheduling on-line advertising, autonomous navigation, etc.

  We will consider some case studies as a part of the course.

- Previous years’ course pages have weekly schedule, readings, references, assignments, assessments. Take a look.

Happy decision making!